

```
In [1]: import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras import layers, models, preprocessing
from sklearn.preprocessing import LabelEncoder
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
import pickle

np.random.seed(1234)
```

Question 1

Reading in file from Elon Musk's Tweets Dataset 2022 from Kaggle.com

```
In [2]: df = pd.read_csv('emails.csv')
df.shape
```

Out[2]: (5172, 3002)

```
In [3]: df.head()
```

```
Out[3]:
```

	Email No.	the	to	ect	and	for	of	a	you	hou	...	connevey	jay	valued	lay	infrastructure	military	allowing	ff	dry	Prediction
0	Email 1	0	0	1	0	0	0	2	0	0	...	0	0	0	0	0	0	0	0	0	0
1	Email 2	8	13	24	6	6	2	102	1	27	...	0	0	0	0	0	0	0	1	0	0
2	Email 3	0	0	1	0	0	0	8	0	0	...	0	0	0	0	0	0	0	0	0	0
3	Email 4	0	5	22	0	5	1	51	2	10	...	0	0	0	0	0	0	0	0	0	0
4	Email 5	7	6	17	1	5	2	57	0	9	...	0	0	0	0	0	0	0	1	0	0

5 rows × 3002 columns

Divide into Train/Test

```
In [4]: data = df.values
X, y = data[:, 1:3001], data[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size = 0.2, random_state=1)

print(X.shape, y.shape)
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
X_train = np.asarray(X_train).astype('float32')
X_test = np.asarray(X_test).astype('float32')
y_train = np.asarray(y_train).astype('float32')
y_test = np.asarray(y_test).astype('float32')

(5172, 3000) (5172,)
(4137, 3000) (1035, 3000) (4137,) (1035,)
```

Question 2

Create a sequential model and evaluate on the test data

```
In [5]: from keras.models import Sequential
model = Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(3000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

```
In [6]: X_val = X_train[:3000]
partial_X_train = X_train[3000:]

y_val = y_train[:3000]
partial_y_train = y_train[3000:]
```

```
In [7]: history = model.fit(partial_X_train,
```

```
partial_y_train,  
epochs=20,  
batch_size=512,  
validation_data=(X_val, y_val))
```

```
Epoch 1/20  
3/3 [=====] - 1s 93ms/step - loss: 1.6857 - accuracy: 0.3096 - val_loss: 0.7017 - val_accuracy: 0.6220  
Epoch 2/20  
3/3 [=====] - 0s 26ms/step - loss: 0.7132 - accuracy: 0.6297 - val_loss: 0.7041 - val_accuracy: 0.7260  
Epoch 3/20  
3/3 [=====] - 0s 24ms/step - loss: 0.6711 - accuracy: 0.7247 - val_loss: 0.5921 - val_accuracy: 0.7720  
Epoch 4/20  
3/3 [=====] - 0s 27ms/step - loss: 0.5825 - accuracy: 0.7995 - val_loss: 0.6026 - val_accuracy: 0.7720  
Epoch 5/20  
3/3 [=====] - 0s 25ms/step - loss: 0.6078 - accuracy: 0.7757 - val_loss: 0.5739 - val_accuracy: 0.7863  
Epoch 6/20  
3/3 [=====] - 0s 25ms/step - loss: 0.5737 - accuracy: 0.7995 - val_loss: 0.5225 - val_accuracy: 0.8490  
Epoch 7/20  
3/3 [=====] - 0s 25ms/step - loss: 0.5055 - accuracy: 0.8584 - val_loss: 0.4722 - val_accuracy: 0.8753  
Epoch 8/20  
3/3 [=====] - 0s 24ms/step - loss: 0.4618 - accuracy: 0.8883 - val_loss: 0.4404 - val_accuracy: 0.8880  
Epoch 9/20  
3/3 [=====] - 0s 26ms/step - loss: 0.4379 - accuracy: 0.8997 - val_loss: 0.5369 - val_accuracy: 0.7957  
Epoch 10/20  
3/3 [=====] - 0s 25ms/step - loss: 0.4803 - accuracy: 0.8470 - val_loss: 0.4018 - val_accuracy: 0.9047  
Epoch 11/20  
3/3 [=====] - 0s 25ms/step - loss: 0.3953 - accuracy: 0.9120 - val_loss: 0.3818 - val_accuracy: 0.9077  
Epoch 12/20  
3/3 [=====] - 0s 25ms/step - loss: 0.3721 - accuracy: 0.9270 - val_loss: 0.3614 - val_accuracy: 0.9037  
Epoch 13/20  
3/3 [=====] - 0s 25ms/step - loss: 0.3505 - accuracy: 0.9323 - val_loss: 0.3460 - val_accuracy: 0.9150  
Epoch 14/20  
3/3 [=====] - 0s 25ms/step - loss: 0.3367 - accuracy: 0.9305 - val_loss: 0.3269 - val_accuracy: 0.9160  
Epoch 15/20  
3/3 [=====] - 0s 24ms/step - loss: 0.3118 - accuracy: 0.9384 - val_loss: 0.3433 - val_accuracy: 0.8787  
Epoch 16/20  
3/3 [=====] - 0s 24ms/step - loss: 0.4183 - accuracy: 0.9050 - val_loss: 0.3136 - val_accuracy: 0.9303  
Epoch 17/20  
3/3 [=====] - 0s 25ms/step - loss: 0.2898 - accuracy: 0.9455 - val_loss: 0.3069 - val_accuracy: 0.9333  
Epoch 18/20  
3/3 [=====] - 0s 23ms/step - loss: 0.2844 - accuracy: 0.9402 - val_loss: 0.2816 - val_accuracy: 0.9290  
Epoch 19/20  
3/3 [=====] - 0s 23ms/step - loss: 0.2625 - accuracy: 0.9604 - val_loss: 0.2818 - val_accuracy: 0.9407
```

Epoch 20/20

3/3 [=====] - 0s 23ms/step - loss: 0.2785 - accuracy: 0.9402 - val_loss: 0.3180 - val_accuracy: 0.8823

```
In [8]: from sklearn.metrics import classification_report
```

```
pred = model.predict(X_test)
pred = [1.0 if p >= 0.5 else 0.0 for p in pred]
print(classification_report(y_test, pred))
```

```
33/33 [=====] - 0s 1ms/step
```

	precision	recall	f1-score	support
0.0	0.85	1.00	0.92	719
1.0	0.99	0.60	0.75	316
accuracy			0.88	1035
macro avg	0.92	0.80	0.83	1035
weighted avg	0.89	0.88	0.86	1035

```
In [9]: losses_and_metrics = model.evaluate(X_test, y_test, batch_size=128)
losses_and_metrics
```

9/9 [=====] - 0s 2ms/step - loss: 0.3538 - accuracy: 0.8754

```
Out[9]: [0.3537578880786896, 0.8753623366355896]
```

```
In [10]: classes = model.predict(X_test, batch_size=128)
classes[:5]
```

```
9/9 [=====] - 0s 2ms/step
```

```
Out[10]: array([[0.01126047],
 [0.7040208 ],
 [0.03501591],
 [0.11170562],
 [0.99775326]], dtype=float32)
```

```
In [11]: import matplotlib.pyplot as plt
```

```
loss = history.history['loss']
val_loss = history.history['val_loss']
```


```

epochs = range(1, len(loss)+1)

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()

```



In [12]:

```


plt.clf() # clear

acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()

```



```
In [13]: from tensorflow.keras.callbacks import EarlyStopping
```

```
early_stopping = EarlyStopping(
    min_delta=0.001, # min change in metrics
    patience=4,      # min epochs
    restore_best_weights=True,
)
```

```
In [14]: # fit the model again with early stopping
```

```
model.fit(X_train, y_train, callbacks=[early_stopping], batch_size=512)
results = model.evaluate(X_test, y_test)
```

```
1/9 [==>.....] - ETA: 0s - loss: 0.3126 - accuracy: 0.8965WARNING:tensorflow:Early stopping conditioned on metric `val_loss` which is not available. Available metrics are: loss,accuracy
9/9 [=====] - 0s 4ms/step - loss: 0.2656 - accuracy: 0.9355
33/33 [=====] - 0s 1ms/step - loss: 0.7090 - accuracy: 0.7643
```

```
In [15]: results
```

```
Out[15]: [0.7089790105819702, 0.7642512321472168]
```

```
In [16]: pred = model.predict(X_test)
pred = [1.0 if p>= 0.5 else 0.0 for p in pred]
print(classification_report(y_test, pred))
```

```
33/33 [=====] - 0s 1ms/step
              precision    recall  f1-score   support

         0.0         0.75         1.00         0.85         719
         1.0         1.00         0.23         0.37         316

 accuracy          0.87
 macro avg         0.82
 weighted avg      0.82
```

Question 3

Try a different architecture like RNN, CNN, etc and evaluate on the test data

```
In [17]: #RNN
max_features = 10000
maxlen = 500
batch_size = 32

model = models.Sequential()
model.add(layers.Embedding(max_features, 32))
model.add(layers.SimpleRNN(32))
model.add(layers.Dense(1, activation='sigmoid'))

model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, None, 32)	320000
simple_rnn (SimpleRNN)	(None, 32)	2080
dense_3 (Dense)	(None, 1)	33

```
=====
Total params: 322,113
Trainable params: 322,113
Non-trainable params: 0
=====
```

```
In [18]: model.compile(optimizer='rmsprop',
                      loss='binary_crossentropy',
                      metrics=['accuracy'])
```

```
In [19]: history = model.fit(X_train,
                             y_train,
                             epochs=10,
                             batch_size=128,
                             validation_split=0.2)
```

```
Epoch 1/10
26/26 [=====] - 27s 1s/step - loss: 0.6008 - accuracy: 0.7123 - val_loss: 0.5873 - val_accuracy: 0.7186
Epoch 2/10
26/26 [=====] - 27s 1s/step - loss: 0.5900 - accuracy: 0.7126 - val_loss: 0.5742 - val_accuracy: 0.7186
Epoch 3/10
26/26 [=====] - 27s 1s/step - loss: 0.5830 - accuracy: 0.7153 - val_loss: 0.5654 - val_accuracy: 0.7162
Epoch 4/10
26/26 [=====] - 26s 1s/step - loss: 0.5844 - accuracy: 0.7147 - val_loss: 0.5679 - val_accuracy: 0.7150
Epoch 5/10
26/26 [=====] - 26s 1s/step - loss: 0.5768 - accuracy: 0.7202 - val_loss: 0.5656 - val_accuracy: 0.7150
Epoch 6/10
26/26 [=====] - 26s 1s/step - loss: 0.5780 - accuracy: 0.7153 - val_loss: 0.5694 - val_accuracy: 0.7138
Epoch 7/10
26/26 [=====] - 26s 1s/step - loss: 0.5831 - accuracy: 0.7126 - val_loss: 0.5725 - val_accuracy: 0.7186
Epoch 8/10
26/26 [=====] - 26s 1s/step - loss: 0.5761 - accuracy: 0.7186 - val_loss: 0.5590 - val_accuracy: 0.7162
Epoch 9/10
26/26 [=====] - 26s 1s/step - loss: 0.5797 - accuracy: 0.7183 - val_loss: 0.5663 - val_accuracy: 0.7138
Epoch 10/10
26/26 [=====] - 26s 1s/step - loss: 0.5732 - accuracy: 0.7220 - val_loss: 0.5672 - val_accuracy: 0.7138
```

```
In [21]: pred = model.predict(X_test)
pred = [1.0 if p >= 0.5 else 0.0 for p in pred]
print(classification_report(y_test, pred))
```



```

33/33 [=====] - 3s 78ms/step
          precision    recall  f1-score   support

     0.0         0.70      0.99      0.82         719
     1.0         0.68      0.05      0.09         316

 accuracy                   0.70         1035
 macro avg              0.69      0.52      0.46         1035
 weighted avg          0.70      0.70      0.60         1035

```

In [22]:

```

#LSTM
model = models.Sequential()
model.add(layers.Embedding(max_features, 32))
model.add(layers.LSTM(32))
model.add(layers.Dense(1, activation='sigmoid'))

model.summary()

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, None, 32)	320000
lstm (LSTM)	(None, 32)	8320
dense_4 (Dense)	(None, 1)	33

=====
 Total params: 328,353
 Trainable params: 328,353
 Non-trainable params: 0

In [23]:

```

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])

```

In [24]:

```

history = model.fit(X_train,
                    y_train,

```

```
epochs=10,
batch_size=128,
validation_split=0.2)
```

```
Epoch 1/10
26/26 [=====] - 65s 2s/step - loss: 0.6266 - accuracy: 0.6951 - val_loss: 0.5935 - val_accuracy: 0.7186
Epoch 2/10
26/26 [=====] - 64s 2s/step - loss: 0.5967 - accuracy: 0.7126 - val_loss: 0.5919 - val_accuracy: 0.7186
Epoch 3/10
26/26 [=====] - 63s 2s/step - loss: 0.5961 - accuracy: 0.7126 - val_loss: 0.5880 - val_accuracy: 0.7186
Epoch 4/10
26/26 [=====] - 63s 2s/step - loss: 0.5933 - accuracy: 0.7126 - val_loss: 0.5880 - val_accuracy: 0.7186
Epoch 5/10
26/26 [=====] - 63s 2s/step - loss: 0.5905 - accuracy: 0.7126 - val_loss: 0.5822 - val_accuracy: 0.7186
Epoch 6/10
26/26 [=====] - 64s 2s/step - loss: 0.5889 - accuracy: 0.7123 - val_loss: 0.5779 - val_accuracy: 0.7186
Epoch 7/10
26/26 [=====] - 64s 2s/step - loss: 0.5864 - accuracy: 0.7126 - val_loss: 0.5756 - val_accuracy: 0.7198
Epoch 8/10
26/26 [=====] - 63s 2s/step - loss: 0.5841 - accuracy: 0.7123 - val_loss: 0.5734 - val_accuracy: 0.7198
Epoch 9/10
26/26 [=====] - 63s 2s/step - loss: 0.5824 - accuracy: 0.7135 - val_loss: 0.5721 - val_accuracy: 0.7186
Epoch 10/10
26/26 [=====] - 63s 2s/step - loss: 0.5826 - accuracy: 0.7129 - val_loss: 0.5700 - val_accuracy: 0.7186
```

In [25]:

```
pred = model.predict(X_test)
pred = [1.0 if p>= 0.5 else 0.0 for p in pred]
print(classification_report(y_test, pred))
```

```
33/33 [=====] - 7s 195ms/step
              precision    recall  f1-score   support

     0.0         0.70      0.99      0.82         719
     1.0         0.42      0.02      0.03         316

 accuracy          0.69         1035
 macro avg         0.56         0.50         0.42         1035
 weighted avg         0.61         0.69         0.58         1035
```

In [26]:

```
#GRU
model = models.Sequential()
```

```

model.add(layers.Embedding(max_features, 32))
model.add(layers.GRU(32))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])

history = model.fit(X_train,
                    y_train,
                    epochs=10,
                    batch_size=128,
                    validation_split=0.2)

```

```

Epoch 1/10
26/26 [=====] - 78s 3s/step - loss: 0.6206 - accuracy: 0.6948 - val_loss: 0.5913 - val_accuracy: 0.7186
Epoch 2/10
26/26 [=====] - 70s 3s/step - loss: 0.5940 - accuracy: 0.7126 - val_loss: 0.5838 - val_accuracy: 0.7186
Epoch 3/10
26/26 [=====] - 69s 3s/step - loss: 0.5908 - accuracy: 0.7126 - val_loss: 0.5808 - val_accuracy: 0.7186
Epoch 4/10
26/26 [=====] - 69s 3s/step - loss: 0.5893 - accuracy: 0.7123 - val_loss: 0.5776 - val_accuracy: 0.7186
Epoch 5/10
26/26 [=====] - 69s 3s/step - loss: 0.5869 - accuracy: 0.7123 - val_loss: 0.5763 - val_accuracy: 0.7162
Epoch 6/10
26/26 [=====] - 67s 3s/step - loss: 0.5856 - accuracy: 0.7120 - val_loss: 0.5730 - val_accuracy: 0.7174
Epoch 7/10
26/26 [=====] - 67s 3s/step - loss: 0.5833 - accuracy: 0.7120 - val_loss: 0.5730 - val_accuracy: 0.7174
Epoch 8/10
26/26 [=====] - 66s 3s/step - loss: 0.5824 - accuracy: 0.7117 - val_loss: 0.5698 - val_accuracy: 0.7186
Epoch 9/10
26/26 [=====] - 67s 3s/step - loss: 0.5814 - accuracy: 0.7120 - val_loss: 0.5689 - val_accuracy: 0.7186
Epoch 10/10
26/26 [=====] - 67s 3s/step - loss: 0.5819 - accuracy: 0.7123 - val_loss: 0.5691 - val_accuracy: 0.7186

```

In [28]:

```

pred = model.predict(X_test)
pred = [1.0 if p >= 0.5 else 0.0 for p in pred]
print(classification_report(y_test, pred))

```

```

33/33 [=====] - 6s 162ms/step
              precision    recall  f1-score   support

0.0          0.69         0.99         0.82         719

```

	1.0	0.27	0.01	0.02	316
accuracy				0.69	1035
macro avg	0.48	0.50	0.42		1035
weighted avg	0.57	0.69	0.57		1035

In [29]:

```
#CNN
max_features = 3000
maxlen = 3000
batch_size = 32

model = models.Sequential()
model.add(layers.Embedding(max_features, 128, input_length=maxlen))
model.add(layers.Conv1D(32, 7, activation='relu'))
model.add(layers.MaxPooling1D(5))
model.add(layers.Conv1D(32, 7, activation='relu'))
model.add(layers.GlobalMaxPooling1D())
model.add(layers.Dense(1))
```

In [30]:

```
model.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, 3000, 128)	384000
conv1d (Conv1D)	(None, 2994, 32)	28704
max_pooling1d (MaxPooling1D)	(None, 598, 32)	0
conv1d_1 (Conv1D)	(None, 592, 32)	7200
global_max_pooling1d (GlobalMaxPooling1D)	(None, 32)	0
dense_6 (Dense)	(None, 1)	33

=====
Total params: 419,937

Trainable params: 419,937

Non-trainable params: 0

In [31]:

```
model.compile(optimizer=tf.keras.optimizers.RMSprop(learning_rate=1e-4), # set learning rate
              loss='binary_crossentropy',
              metrics=['accuracy'])

history = model.fit(X_train,
                    y_train,
                    epochs=10,
                    batch_size=128,
                    validation_split=0.2)
```

Epoch 1/10

26/26 [=====] - 23s 856ms/step - loss: 0.6220 - accuracy: 0.7126 - val_loss: 0.5919 - val_accuracy: 0.7186

Epoch 2/10

26/26 [=====] - 22s 843ms/step - loss: 0.5944 - accuracy: 0.7126 - val_loss: 0.5882 - val_accuracy: 0.7186

Epoch 3/10

26/26 [=====] - 22s 843ms/step - loss: 0.5877 - accuracy: 0.7126 - val_loss: 0.5854 - val_accuracy: 0.7186

Epoch 4/10

26/26 [=====] - 22s 850ms/step - loss: 0.5820 - accuracy: 0.7126 - val_loss: 0.5821 - val_accuracy: 0.7186

Epoch 5/10

26/26 [=====] - 22s 845ms/step - loss: 0.5769 - accuracy: 0.7126 - val_loss: 0.5763 - val_accuracy: 0.7186

Epoch 6/10

26/26 [=====] - 23s 868ms/step - loss: 0.5726 - accuracy: 0.7126 - val_loss: 0.5734 - val_accuracy: 0.7186

Epoch 7/10

26/26 [=====] - 24s 923ms/step - loss: 0.5659 - accuracy: 0.7126 - val_loss: 0.5687 - val_accuracy: 0.7186

Epoch 8/10

26/26 [=====] - 23s 893ms/step - loss: 0.5608 - accuracy: 0.7126 - val_loss: 0.5651 - val_accuracy: 0.7186

Epoch 9/10

26/26 [=====] - 27s 1s/step - loss: 0.5551 - accuracy: 0.7126 - val_loss: 0.5602 - val_accuracy: 0.7186

Epoch 10/10

26/26 [=====] - 24s 930ms/step - loss: 0.5488 - accuracy: 0.7126 - val_loss: 0.5563 - val_accuracy: 0.7186

```
In [32]: from sklearn.metrics import classification_report
pred = model.predict(X_test)
pred = [1.0 if p >= 0.5 else 0.0 for p in pred]
print(classification_report(y_test, pred))
```

```
33/33 [=====] - 1s 36ms/step
      precision    recall  f1-score   support

     0.0         0.69      1.00      0.82        719
     1.0         0.00      0.00      0.00        316

 accuracy          0.69        1035
 macro avg         0.35      0.50      0.41        1035
 weighted avg      0.48      0.69      0.57        1035
```

C:\Users\aurol\anaconda3\lib\site-packages\sklearn\metrics_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

C:\Users\aurol\anaconda3\lib\site-packages\sklearn\metrics_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

C:\Users\aurol\anaconda3\lib\site-packages\sklearn\metrics_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

Question 4

Try different embedding approaches and evaluate on the test data

```
In [33]: max_features = 3000
maxlen = 3000
model = models.Sequential()
model.add(layers.Embedding(max_features, 8, input_length=maxlen))
model.add(layers.Flatten())
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()
```

```
history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_split=0.2)
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, 3000, 8)	24000
flatten (Flatten)	(None, 24000)	0
dense_7 (Dense)	(None, 16)	384016
dense_8 (Dense)	(None, 1)	17

```
=====  
Total params: 408,033  
Trainable params: 408,033  
Non-trainable params: 0
```

```
Epoch 1/10  
104/104 [=====] - 1s 9ms/step - loss: 0.4887 - acc: 0.7610 - val_loss: 0.3561 - val_acc: 0.7717  
Epoch 2/10  
104/104 [=====] - 1s 8ms/step - loss: 0.1920 - acc: 0.9326 - val_loss: 0.1379 - val_acc: 0.9601  
Epoch 3/10  
104/104 [=====] - 1s 8ms/step - loss: 0.0902 - acc: 0.9710 - val_loss: 0.1141 - val_acc: 0.9601  
Epoch 4/10  
104/104 [=====] - 1s 8ms/step - loss: 0.0535 - acc: 0.9837 - val_loss: 0.1000 - val_acc: 0.9614  
Epoch 5/10  
104/104 [=====] - 1s 8ms/step - loss: 0.0358 - acc: 0.9894 - val_loss: 0.0830 - val_acc: 0.9722  
Epoch 6/10  
104/104 [=====] - 1s 8ms/step - loss: 0.0228 - acc: 0.9934 - val_loss: 0.0851 - val_acc: 0.9722  
Epoch 7/10  
104/104 [=====] - 1s 8ms/step - loss: 0.0155 - acc: 0.9952 - val_loss: 0.1177 - val_acc: 0.9650  
Epoch 8/10  
104/104 [=====] - 1s 8ms/step - loss: 0.0101 - acc: 0.9982 - val_loss: 0.1070 - val_acc: 0.9734  
Epoch 9/10  
104/104 [=====] - 1s 8ms/step - loss: 0.0074 - acc: 0.9985 - val_loss: 0.1152 - val_acc: 0.9710  
Epoch 10/10  
104/104 [=====] - 1s 8ms/step - loss: 0.0053 - acc: 0.9988 - val_loss: 0.1305 - val_acc: 0.9710
```

In [34]:

```
pred = model.predict(X_test)  
pred = [1.0 if p>= 0.5 else 0.0 for p in pred]
```

```
print(classification_report(y_test, pred))
```

```
33/33 [=====] - 0s 3ms/step
          precision    recall  f1-score   support

     0.0      0.98      0.98      0.98        719
     1.0      0.96      0.96      0.96        316

 accuracy          0.97          1035
 macro avg          0.97          1035
weighted avg          0.97          1035
```

Question 5

Write up your analysis of the performance of various approaches

From a quick glance the accuracies of the models are listed from lowest to highest: (CNN, GRU, and LSTM) all had .69, RNN with .70, 1st sequential with early stopping at .76, then the first seq without early stopping at .88, and finally embedding method at .97.

Overall the embedding method had the best scores. Followed by 1st seq, 1st seq w/early stopping, RNN, LSTM, GRU, and lastly CNN.